

Data Prep

```
In [5]: print(df.describe())
print('-----')
print(df.Interests.value_counts())
print('-----')
print(df.Gender.value_counts())
print('-----')
print(df.shape)
print(df.isnull().sum())
print(df.info())
```

```

      Unnamed: 0      Age  Annual Income  Total Spending
count  1600.00000  1598.00000    1596.00000    1600.00000
mean    799.50000   37.895494    187.129699    2575.600781
std     462.02453   16.226009   5004.797654    1453.083432
min       0.00000   14.000000   -100.000000   -102.500000
25%     399.75000   28.000000    39.000000    1383.750000
50%     799.50000   35.000000    65.000000    2511.250000
75%    1199.25000   47.000000    81.000000    3843.750000
max    1599.00000  350.000000  200000.000000    5176.250000

```

```

-----
Entertainment      139
Hiking             137
Travel             132
Reading            131
Yoga               116
Crafts             115
Camping            108
Technology          102
Photography         98
Exercise            97
Music              92
Art                88
Gaming             71
Strategic games     57
Sports              37
Pets                28
Automobiles         28
Cooking             19
Gardening           5
Name: Interests, dtype: int64

```

```

-----
female      812
male        788

```

```
Name: Gender, dtype: int64
-----
(1600, 6)
Unnamed: 0      0
Gender          0
Age            2
Annual Income  4
Total Spending  0
Interests       0
dtype: int64
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1600 entries, 0 to 1599
Data columns (total 6 columns):
Unnamed: 0      1600 non-null int64
Gender          1600 non-null object
Age            1598 non-null float64
Annual Income  1596 non-null float64
Total Spending  1600 non-null float64
Interests       1600 non-null object
dtypes: float64(3), int64(1), object(2)
memory usage: 75.1+ KB
None
```

```
In [6]: #Remove data errors, negative values, zeros
df['Annual Income'] = df['Annual Income'].mask(df['Annual Income'] == 200000, 0)
df['Annual Income'] = df['Annual Income'].mask(df['Annual Income'] <= 0, 0)
df['Total Spending'] = df['Total Spending'].mask(df['Total Spending'] <= 0, 0)
df['Age'] = df['Age'].mask(df['Age'] >= 100, 0)
df = remove_nan_and_zeroes_from_columns(df, 'Age')
df = remove_nan_and_zeroes_from_columns(df, 'Annual Income')
df = remove_nan_and_zeroes_from_columns(df, 'Total Spending')
df = df.drop(['Unnamed: 0'],axis = 1)
df.describe()
```

Out[6]:

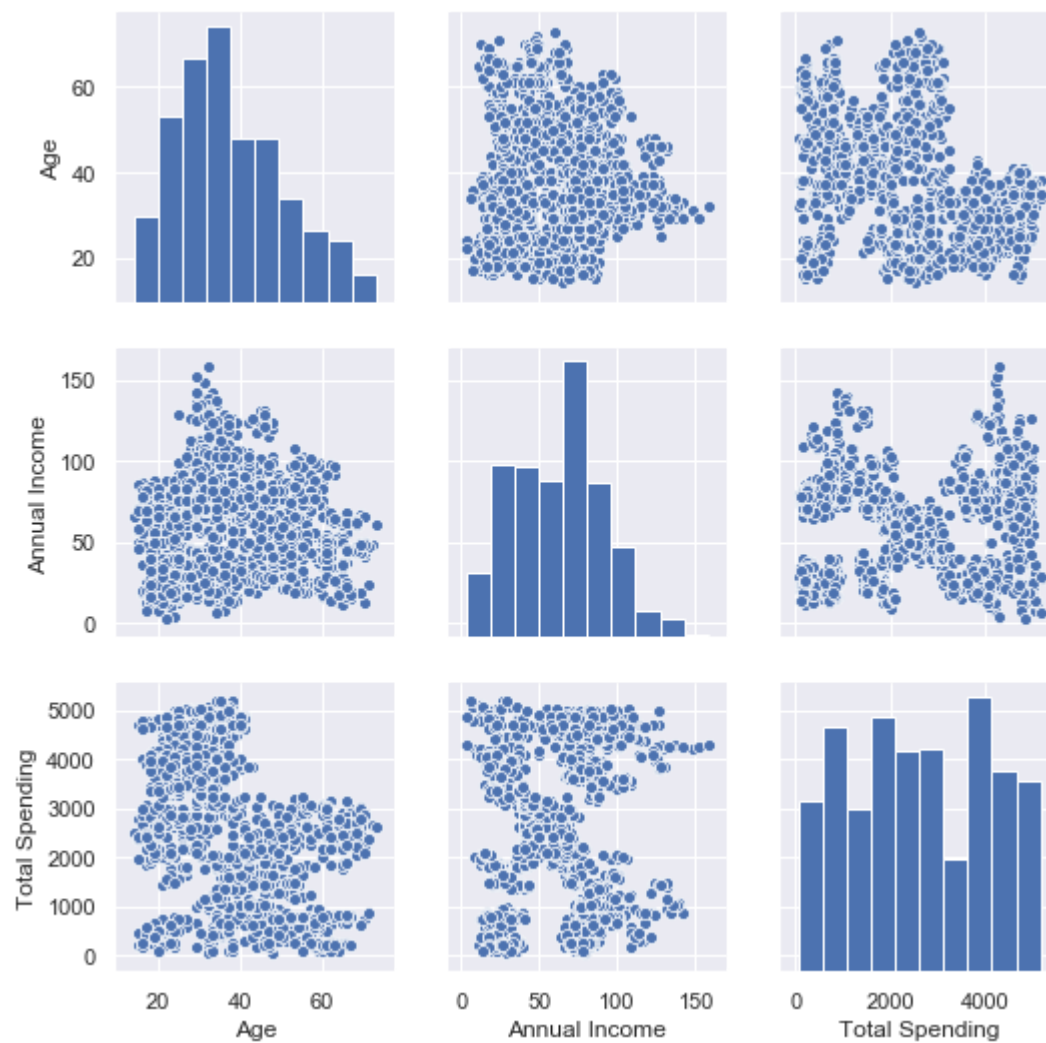
	Age	Annual Income	Total Spending
count	1583.000000	1583.000000	1583.000000
mean	37.569172	61.943146	2585.615919
std	13.232153	28.755595	1446.813956
min	14.000000	3.000000	51.250000
25%	28.000000	39.000000	1435.000000
50%	35.000000	65.000000	2562.500000
75%	47.000000	81.000000	3843.750000
max	73.000000	158.000000	5176.250000

```
In [7]: dfcat = df[['Gender', 'Interests']]
dfnum = df.drop(['Gender', 'Interests'], axis = 1)
```

EDA

```
In [8]: # Basic correlogram  
sns.pairplot(df)
```


```
Out[8]: <seaborn.axisgrid.PairGrid at 0x263885c8400>
```





```
In [9]: # Use the 'hue' argument to provide a factor variable
# sns.set(rc={'figure.figsize':(24,15)})
print(sns.lmplot(x='Total Spending', y="Annual Income", data=df, fit_reg=False, hue='Interests', legend=True))
print(sns.lmplot(x='Total Spending', y="Annual Income", data=df, fit_reg=False, hue='Gender', legend=True))
#sns.plt.show()
```

<seaborn.axisgrid.FacetGrid object at 0x0000026388DACA20>

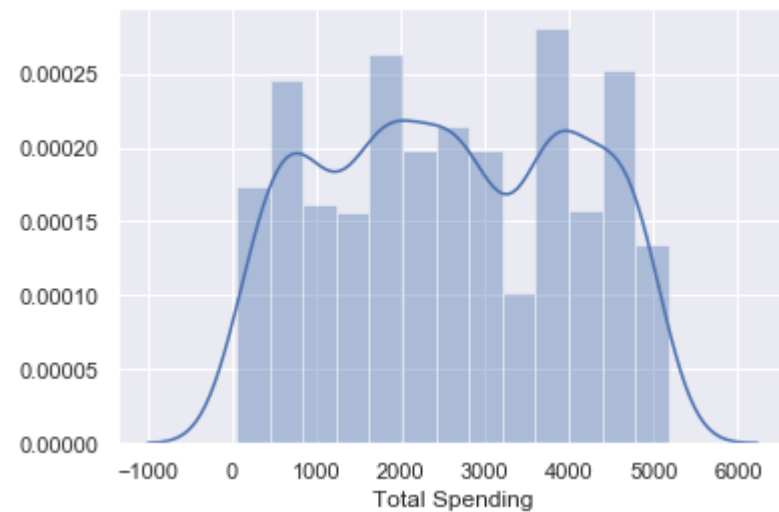
<seaborn.axisgrid.FacetGrid object at 0x0000026388BFB2E8>






```
In [10]: print(sns.distplot(df['Total Spending']))
```

```
AxesSubplot(0.125,0.125;0.775x0.755)
```



```
In [11]: print(sns.distplot(df['Annual Income']))
```

AxesSubplot(0.125,0.125;0.775x0.755)



```
In [12]: print(sns.distplot(df['Age']))
```

AxesSubplot(0.125,0.125;0.775x0.755)



More Data Prep

```
In [13]: df_new = df.copy()
df_new['age_bin'] = pd.cut(df_new['Age'], [0, 20, 30, 40, 50, 60, 70, 80],
                           labels=['0-20', '20-30', '30-40', '40-50', '50-60', '60-70', '70-80'])
df_new = df_new.drop('Age', axis = 1)
```

```
In [14]: df_new['income_bin'] = pd.cut(df_new['Annual Income'], [0, 20, 30, 40, 50, 60, 70, 80, 90, 100, 110, 120, 130, 140, 150, 160],
                                       labels=['0-20', '20-30', '30-40', '40-50', '50-60', '60-70', '70-80', '80-90', '90-100', '100-110', '110-120', '120-130', '130-140', '140-150', '150-160'])
df_new = df_new.drop('Annual Income', axis = 1)
```

```
In [15]: df_new['spending_bin'] = pd.cut(df_new['Total Spending'], [0, 1000, 2000, 3000, 4000, 5000, 6000],
                                         labels=['0-1000', '1000-2000', '2000-3000', '3000-4000', '4000-5000', '5000-6000'])
df_new = df_new.drop('Total Spending', axis = 1)
```

```
In [16]: df_new2 = df_new.copy()

from sklearn import preprocessing
le = preprocessing.LabelEncoder()
df_new2 = df_new.apply(le.fit_transform)
df_new2.head(1000)
```

Out[16]:

	Gender	Interests	age_bin	income_bin	spending_bin
0	1	3	1	0	0
1	1	6	2	0	0
2	1	17	5	0	0
3	1	4	5	0	0
4	0	17	4	7	0
6	1	17	3	7	0
7	0	17	4	7	0
8	0	13	5	7	0
9	0	13	4	8	0
10	0	15	4	7	0
11	1	15	3	8	0

```
In [17]: km_cao = KModes(n_clusters=2, init = "Cao", n_init = 1, verbose=1)
fitClusters_cao = km_cao.fit_predict(df_new2)
fitClusters_cao
```

```
Init: initializing centroids
Init: initializing clusters
Starting iterations...
Run 1, iteration: 1/100, moves: 0, cost: 4967.0
```

Out[17]: array([1, 0, 1, ..., 0, 0, 1], dtype=uint16)

```
In [18]: clusterCentroidsDf = pd.DataFrame(km_cao.cluster_centroids_)
clusterCentroidsDf.columns = df_new2.columns

# Mode of the clusters
clusterCentroidsDf
```

Out[18]:

	Gender	Interests	age_bin	income_bin	spending_bin
0	0	5	2	12	3
1	1	9	1	11	2

```
In [19]: km_huang = KModes(n_clusters=2, init = "Huang", n_init = 1, verbose=1)
fitClusters_huang = km_huang.fit_predict(df_new2)
fitClusters_huang
```

```
Init: initializing centroids
Init: initializing clusters
Starting iterations...
Run 1, iteration: 1/100, moves: 156, cost: 5066.0
Run 1, iteration: 2/100, moves: 0, cost: 5066.0
```

Out[19]: array([0, 0, 0, ..., 0, 0, 0], dtype=uint16)

```
In [20]: cost = []

for num_clusters in list(range(1,10)):
    kmode = KModes(n_clusters=num_clusters, init = "Cao", n_init = 1, verbose=1)
    kmode.fit_predict(df_new2)
    cost.append(kmode.cost_)

y = np.array([i for i in range(1,10,1)])
plt.plot(y,cost)
```


```
Init: initializing centroids
Init: initializing clusters
Starting iterations...
Run 1, iteration: 1/100, moves: 0, cost: 5884.0
Init: initializing centroids
Init: initializing clusters
Starting iterations...
Run 1, iteration: 1/100, moves: 0, cost: 4967.0
Init: initializing centroids
Init: initializing clusters
Starting iterations...
Run 1, iteration: 1/100, moves: 0, cost: 4576.0
Init: initializing centroids
Init: initializing clusters
Starting iterations...
Run 1, iteration: 1/100, moves: 0, cost: 4312.0
Init: initializing centroids
Init: initializing clusters
Starting iterations...
Run 1, iteration: 1/100, moves: 35, cost: 4068.0
Run 1, iteration: 2/100, moves: 0, cost: 4068.0
Init: initializing centroids
Init: initializing clusters
Starting iterations...
Run 1, iteration: 1/100, moves: 33, cost: 3881.0
Run 1, iteration: 2/100, moves: 0, cost: 3881.0
Init: initializing centroids
Init: initializing clusters
Starting iterations...
Run 1, iteration: 1/100, moves: 29, cost: 3777.0
Run 1, iteration: 2/100, moves: 0, cost: 3777.0
Init: initializing centroids
Init: initializing clusters
```

```

Starting iterations...
Run 1, iteration: 1/100, moves: 89, cost: 3677.0
Run 1, iteration: 2/100, moves: 0, cost: 3677.0
Init: initializing centroids
Init: initializing clusters
Starting iterations...
Run 1, iteration: 1/100, moves: 87, cost: 3603.0
Run 1, iteration: 2/100, moves: 0, cost: 3603.0

```

Out[20]: [<matplotlib.lines.Line2D at 0x2638964c5c0>]



```

In [21]: km_cao = KModes(n_clusters=3, init = "Cao", n_init = 1, verbose=1)
fitClusters_cao = km_cao.fit_predict(df_new2)
fitClusters_cao

```

```

Init: initializing centroids
Init: initializing clusters
Starting iterations...
Run 1, iteration: 1/100, moves: 0, cost: 4576.0

```

Out[21]: array([1, 2, 1, ..., 0, 0, 1], dtype=uint16)

Post-processesing & prediction evaluation

```
In [22]: clustersDf = pd.DataFrame(fitClusters_cao)
clustersDf.columns = ['cluster_predicted']
combinedDf = pd.concat([df_new, clustersDf], axis = 1).reset_index()
combinedDf = combinedDf.drop(['index'], axis = 1)

combinedDf.head()
```

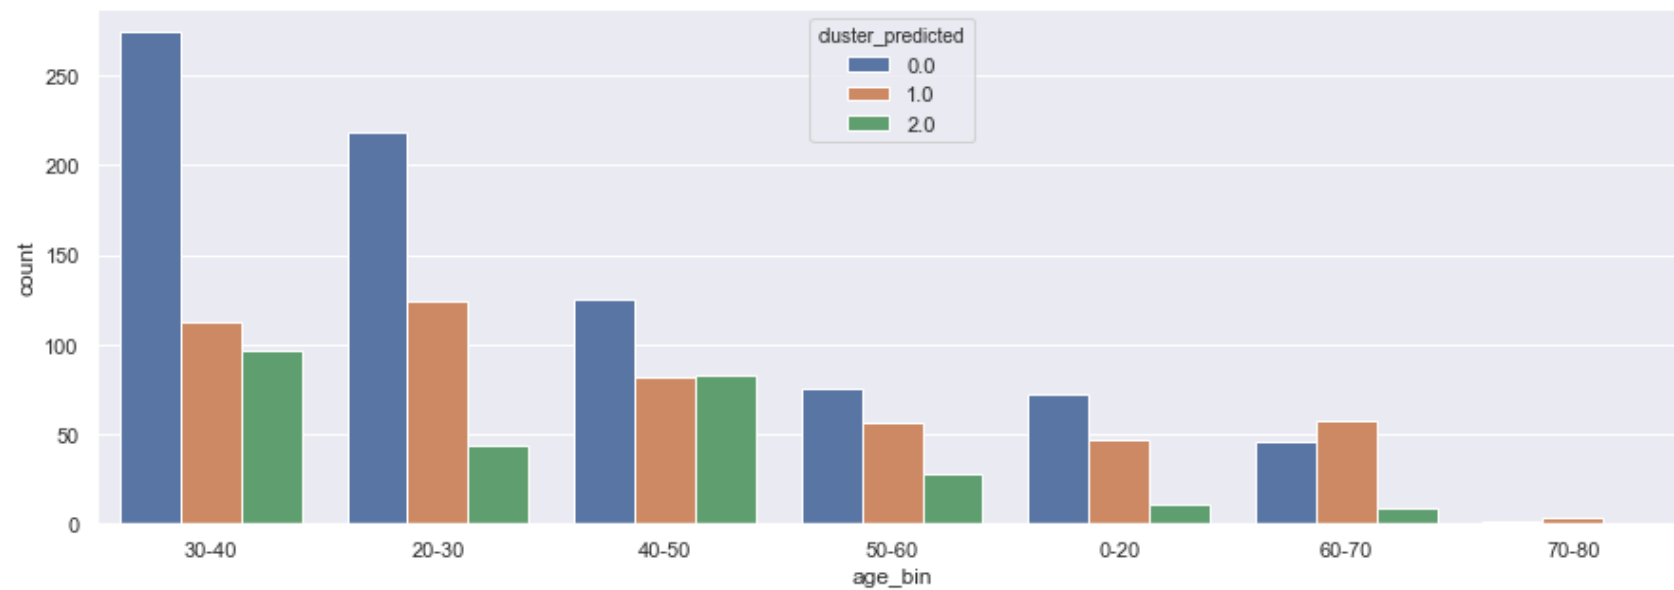
Out[22]:

	Gender	Interests	age_bin	income_bin	spending_bin	cluster_predicted
0	male	Cooking	20-30	0-20	0-1000	1.0
1	male	Exercise	30-40	0-20	0-1000	2.0
2	male	Travel	60-70	0-20	0-1000	1.0
3	male	Crafts	60-70	0-20	0-1000	1.0
4	female	Travel	50-60	20-30	0-1000	0.0


```
In [23]: cluster_0 = combinedDf[combinedDf['cluster_predicted'] == 0]
cluster_1 = combinedDf[combinedDf['cluster_predicted'] == 1]
cluster_2 = combinedDf[combinedDf['cluster_predicted'] == 2]
```



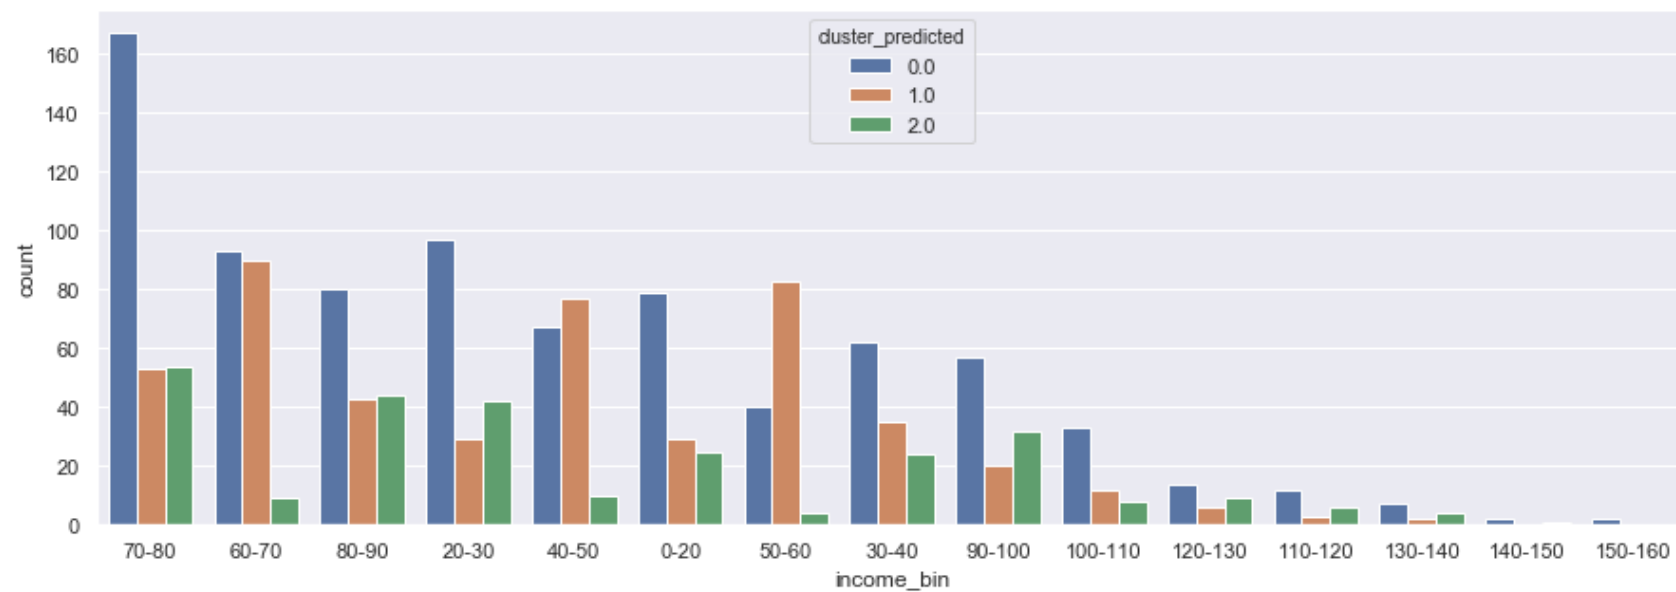
```
In [24]: plt.subplots(figsize = (15,5))  
sns.countplot(x=combinedDf['age_bin'],order=combinedDf['age_bin'].value_counts().index,hue=combinedDf['cluster_p',  
plt.show()
```



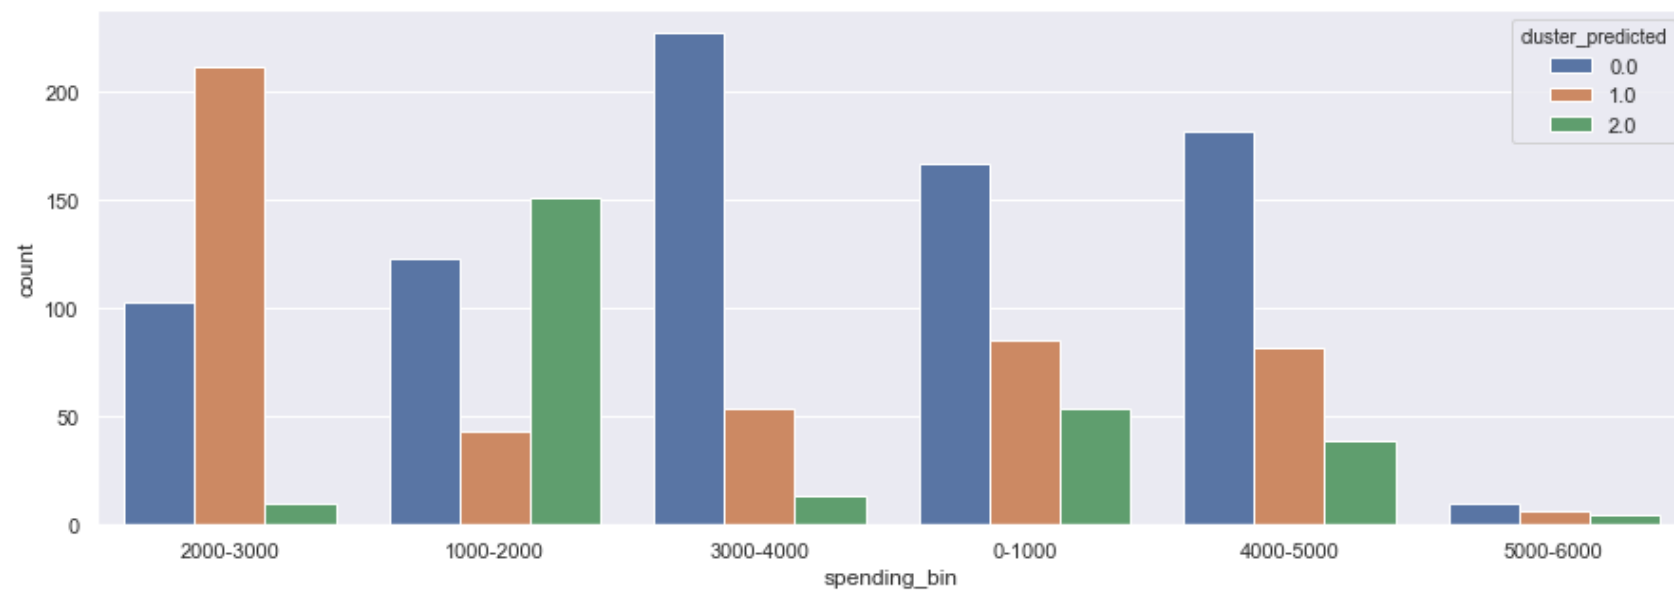
```
In [25]: plt.subplots(figsize = (15,5))
sns.countplot(x=combinedDf['Interests'],order=combinedDf['Interests'].value_counts().index,hue=combinedDf['cluster'],
plt.show()
```



```
In [26]: plt.subplots(figsize = (15,5))  
sns.countplot(x=combinedDf['income_bin'],order=combinedDf['income_bin'].value_counts().index,hue=combinedDf['cluster'],  
plt.show()
```




```
In [27]: plt.subplots(figsize = (15,5))  
sns.countplot(x=combinedDf['spending_bin'],order=combinedDf['spending_bin'].value_counts().index,hue=combinedDf['cluster_predicted'],plt.show()
```



```
In [31]: combinedDf2 = pd.concat([df, clustersDf], axis = 1).reset_index()
combinedDf2 = combinedDf2.drop(['index'], axis = 1)

combinedDf2.head()

print(sns.lmplot(x='Total Spending', y="Annual Income", data=combinedDf2, fit_reg=False, hue='cluster_predicted')
<seaborn.axisgrid.FacetGrid object at 0x0000026389AB56D8>
```



In []: