

Simple Classifiers Exercise

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import nltk
import os

from nltk.tokenize import WordPunctTokenizer
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer

plt.style.use('seaborn-deep')

%matplotlib inline
# Eliminates output truncation
pd.options.display.max_columns = 999
# Use seaborn style defaults and set the default figure size
pd.set_option('display.max_rows', None)
```

```
In [2]: os.chdir('C:\\Users\\orion.darley\\Desktop\\ML HW\\')
cwd = os.getcwd()
```

Load

```
In [3]: url = 'https://raw.githubusercontent.com/OrionDarley/Public-Other/master/News%20Story%20NLP%20Classifier/train'
train = pd.read_csv(url, error_bad_lines=False).rename(columns={'Content': 'content', 'Label': 'label'})
url = 'https://raw.githubusercontent.com/OrionDarley/Public-Other/master/News%20Story%20NLP%20Classifier/valid'
valid = pd.read_csv(url, error_bad_lines=False)
```


Data Prep & CV / TFIDF Vectorization and Matrix

Tf-idf is a scoring scheme for words - that is a measure of how important a word is to a document. I chose Tf-idf over GloVe, because GloVe and word2vec are models that learn from vectors of words by taking into consideration their occurrence and co-occurrence information. While word2vec can be seen as a model that improves its ability to predict [(target word | context words), and GloVe is modeled to do dimensionality reduction.]

```
In [4]: y_train = train.label.replace({'Classifieds': 1, 'News': 2, 'Features': 3, 'Opinion': 4})
y_valid = valid.label.replace({'Classifieds': 1, 'News': 2, 'Features': 3, 'Opinion': 4})
x_train = train.content
x_valid = valid.content
```

```
In [5]: bins = np.linspace(0, 6, 10)

plt.hist([y_train, y_valid], bins, label=['train.y', 'valid.y'])
plt.legend(loc='upper right')
plt.show()
```



```
In [6]: #Fit models on tfidf matrix
wpt = WordPunctTokenizer()
cv = CountVectorizer(tokenizer=wpt.tokenize, stop_words='english', ngram_range=(1,3))

cv.fit(train.content.tolist() + valid.content.tolist())
train_cv = cv.transform(x_train)
valid_cv = cv.transform(x_valid)

tfidf = TfidfVectorizer(min_df=3, max_features=None,
                        strip_accents='unicode', analyzer='word', token_pattern=r'\w{1,}',
                        ngram_range=(1, 3), use_idf=1.0, smooth_idf=1.0, sublinear_tf=1.0,
                        stop_words = 'english')

tfidf.fit(x_train.tolist() + x_valid.tolist())
train_tf = tfidf.transform(x_train)
valid_tf = tfidf.transform(x_valid)
print(train_tf.shape, valid_tf.shape)
```

```
(2000, 50198) (2000, 50198)
```

ML Classifiers using TFIDF Matrix

```
In [7]: from sklearn.metrics import accuracy_score, log_loss, classification_report
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC, LinearSVC, NuSVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, GradientBoostingClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis, QuadraticDiscriminantAnalysis
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV, cross_val_score

#Creates confusion matrix
def Model(train_tf, y_train, valid_tf, y_valid, model):
    model.fit(train_tf, y_train)
    prediction = model.predict(valid_tf)
    print(classification_report(y_valid, prediction))

classifiers = [
    LogisticRegression(),
    MultinomialNB(),
    KNeighborsClassifier(3),
    SVC(kernel="rbf", C=0.025, probability=True),
    NuSVC(probability=True),
    DecisionTreeClassifier(),
    RandomForestClassifier(),
    AdaBoostClassifier(),
    GradientBoostingClassifier()]

names = ["LogisticRegression", "MultinomialNB", "KNeighborsClassifier", "SVC",
         "NuSVC", "DecisionTreeClassifier", "RandomForestClassifier", "AdaBoostClassifier",
         "GradientBoostingClassifier"]
```

```
In [8]: df = pd.DataFrame({'unnamed': [0]})
tfidfcores = pd.DataFrame()

for classifier, name in zip(classifiers, names):
    pipe = Pipeline(steps=[('classifier', classifier)])
    pipe.fit(train_tf, y_train)
    #print(classifier)
    print('-----')
    print(name)
    print("model ave weighted f1-score: %.3f" % pipe.score(valid_tf, y_valid))
    Model(train_tf, y_train, valid_tf, y_valid, classifier)
    df['classifier'] = name
    df['score'] = pipe.score(valid_tf, y_valid)
    tfidfcores = tfidfcores.append(df)

tfidfcores = tfidfcores.sort_values('score')
plt.barh(tfidfcores['classifier'], tfidfcores['score'], align='center', alpha=0.5)
plt.yticks(tfidfcores['classifier'])
plt.title('Classification Scores using TFIDF Matrix')

plt.show()
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:460: FutureWarning: Default multi_class will be changed to 'auto' in 0.22. Specify the multi_class option to silence this warning.

"this warning.", FutureWarning)

LogisticRegression

model ave weighted f1-score: 0.678

	precision	recall	f1-score	support
1	0.63	0.70	0.66	470
2	0.76	0.56	0.65	512
3	0.59	0.58	0.59	511
4	0.74	0.87	0.80	507
micro avg	0.68	0.68	0.68	2000

macro avg	0.68	0.68	0.67	2000
weighted avg	0.68	0.68	0.67	2000

MultinomialNB

model ave weighted f1-score: 0.668

	precision	recall	f1-score	support
1	0.60	0.72	0.65	470
2	0.75	0.56	0.64	512
3	0.57	0.59	0.58	511
4	0.79	0.80	0.79	507

micro avg	0.67	0.67	0.67	2000
macro avg	0.68	0.67	0.67	2000
weighted avg	0.68	0.67	0.67	2000

KNeighborsClassifier

model ave weighted f1-score: 0.630

	precision	recall	f1-score	support
1	0.52	0.75	0.62	470
2	0.65	0.64	0.65	512
3	0.64	0.47	0.54	511
4	0.77	0.67	0.71	507

micro avg	0.63	0.63	0.63	2000
macro avg	0.64	0.63	0.63	2000
weighted avg	0.65	0.63	0.63	2000

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\svm\base.py:196: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

SVC

model ave weighted f1-score: 0.235

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1143: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.

'precision', 'predicted', average, warn_for)

	precision	recall	f1-score	support
1	0.23	1.00	0.38	470
2	0.00	0.00	0.00	512
3	0.00	0.00	0.00	511
4	0.00	0.00	0.00	507
micro avg	0.23	0.23	0.23	2000
macro avg	0.06	0.25	0.10	2000
weighted avg	0.06	0.23	0.09	2000

NuSVC

model ave weighted f1-score: 0.620

	precision	recall	f1-score	support
1	0.68	0.57	0.62	470
2	0.88	0.37	0.52	512
3	0.45	0.73	0.56	511
4	0.72	0.81	0.77	507
micro avg	0.62	0.62	0.62	2000
macro avg	0.68	0.62	0.62	2000
weighted avg	0.69	0.62	0.61	2000

DecisionTreeClassifier

model ave weighted f1-score: 0.577

	precision	recall	f1-score	support
1	0.51	0.62	0.56	470
2	0.58	0.52	0.55	512
3	0.46	0.47	0.47	511
4	0.79	0.69	0.74	507
micro avg	0.58	0.58	0.58	2000
macro avg	0.58	0.58	0.58	2000
weighted avg	0.59	0.58	0.58	2000

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:246: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.

"10 in version 0.20 to 100 in 0.22.", FutureWarning)

RandomForestClassifier

model ave weighted f1-score: 0.584

	precision	recall	f1-score	support
1	0.55	0.63	0.58	470
2	0.62	0.61	0.61	512
3	0.48	0.41	0.44	511
4	0.75	0.77	0.76	507
micro avg	0.60	0.60	0.60	2000
macro avg	0.60	0.60	0.60	2000
weighted avg	0.60	0.60	0.60	2000

AdaBoostClassifier

model ave weighted f1-score: 0.467

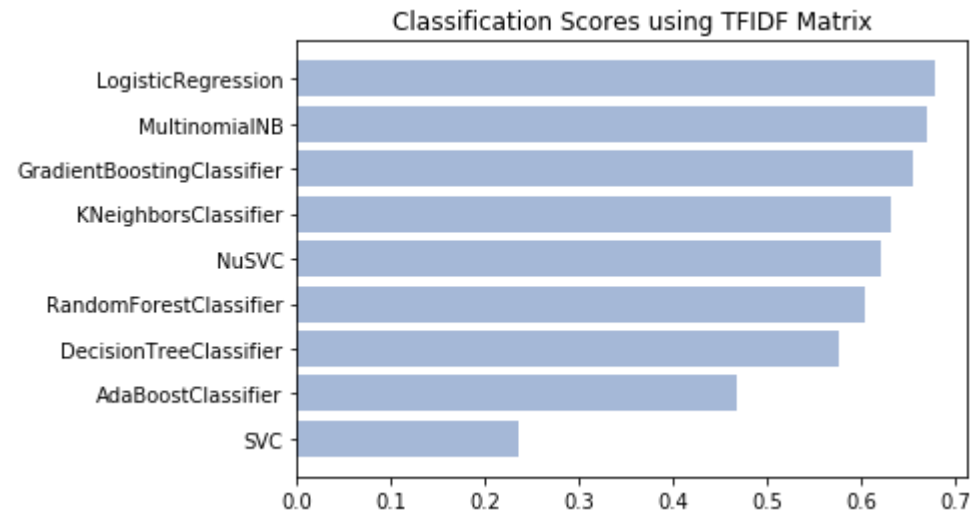
	precision	recall	f1-score	support
1	0.45	0.44	0.45	470
2	0.44	0.54	0.49	512
3	0.30	0.33	0.31	511
4	0.79	0.55	0.65	507
micro avg	0.47	0.47	0.47	2000
macro avg	0.50	0.47	0.47	2000
weighted avg	0.50	0.47	0.47	2000

GradientBoostingClassifier

model ave weighted f1-score: 0.655

	precision	recall	f1-score	support
1	0.57	0.64	0.61	470
2	0.70	0.60	0.64	512
3	0.57	0.54	0.55	511
4	0.78	0.84	0.81	507
micro avg	0.66	0.66	0.66	2000
macro avg	0.65	0.66	0.65	2000

weighted avg 0.66 0.66 0.65 2000



In [9]: `from yellowbrick.classifier import ConfusionMatrix, PrecisionRecallCurve`

```
estimators_dct = {"Logistic Legression": LogisticRegression(),
                  "MultinomialNB" : MultinomialNB(),
                  "GradientBoostingClassifier": GradientBoostingClassifier(),
                  "KNeighborsClassifier": KNeighborsClassifier(3)}

# set up the figure size for the confusion matrices
plt.rcParams['figure.figsize'] = (6, 4)
plt.rcParams['font.size'] = 15

def confusion_matrices(estimator_dct):

    for estimator in estimator_dct.keys():
        print(estimator)
        model = Pipeline([('estimator', estimator_dct[estimator])])
        model.fit(train_tf, y_train)
        cm = ConfusionMatrix(model, fontsize=13)
        cm.score(valid_tf, y_valid)
        cm.poof()

confusion_matrices(estimators_dct)
```

Logistic Legression

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.


FutureWarning)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:460: FutureWarning: Default multi_class will be changed to 'auto' in 0.22. Specify the multi_class option to silence this warning.

"this warning.", FutureWarning)


C:\ProgramData\Anaconda3\lib\site-packages\yellowbrick\classifier\base.py:232: YellowbrickWarning: could not determine class_counts_ from previously fitted classifier

YellowbrickWarning,




MultinomialNB

C:\ProgramData\Anaconda3\lib\site-packages\yellowbrick\classifier\base.py:232: YellowbrickWarning: could not determine class_counts_ from previously fitted classifier
 YellowbrickWarning,



GradientBoostingClassifier



KNeighborsClassifier

```
C:\ProgramData\Anaconda3\lib\site-packages\yellowbrick\classifier\base.py:232: YellowbrickWarning: could not determine class_counts_ from previously fitted classifier
YellowbrickWarning,
```

KNeighborsClassifier Confusion Matrix